

Original decisiontrees Slides



# Decision Trees

CS229: Machine Learning


Carlos Guestrin

Stanford University

Slides include content developed by and co-developed with  
Emily Fox

# Predicting potential loan defaults

# What makes a loan risky?




# Credit history explained

Did I pay previous  
loans on time?




**Example:**  
excellent, good, or  
fair



# Income

What's my income?


**Example:**  
\$80K per year



# Loan terms

How soon do I need to pay the loan?


**Example:** 3 years,  
5 years,...



# Personal information


Age, reason for the loan, marital status,...

**Example:** Home loan for a married couple

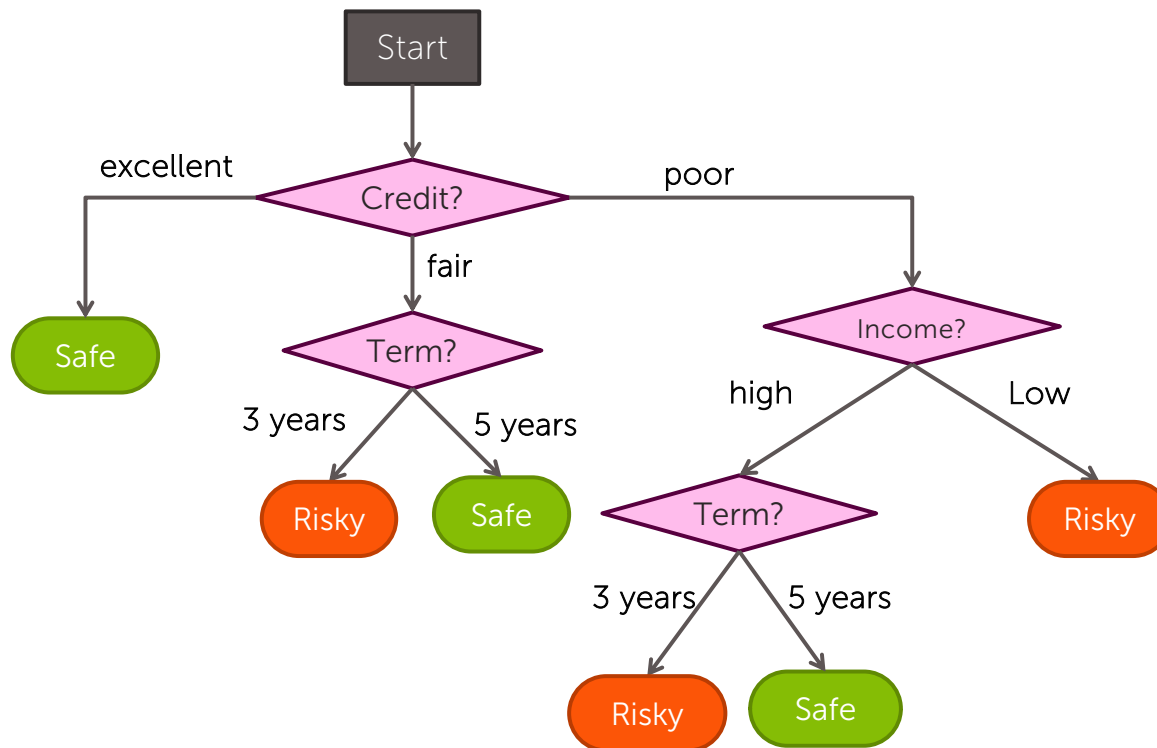





# Classifier review



# This module ... decision trees



# Scoring a loan application




# Decision tree learning task

# Decision tree learning problem

Training data:  $N$  observations  $(\mathbf{x}_i, y_i)$

Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe



# Quality metric: Classification error

- Error measures fraction of mistakes


$$\text{Error} = \frac{\text{\# incorrect predictions}}{\text{\# examples}}$$

- Best possible value : 0.0
- Worst possible value: 1.0

# How do we find the best tree?

Exponentially large number of possible trees makes decision tree learning **hard**!

Learning the smallest decision tree is an *NP-hard problem*  
[Hyafil & Rivest '76]



# Greedy decision tree learning




# Our training data table

Assume  $N = 40$ , 3 features

Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe


# Start with all the data

Loan status: **Safe** **Risky**

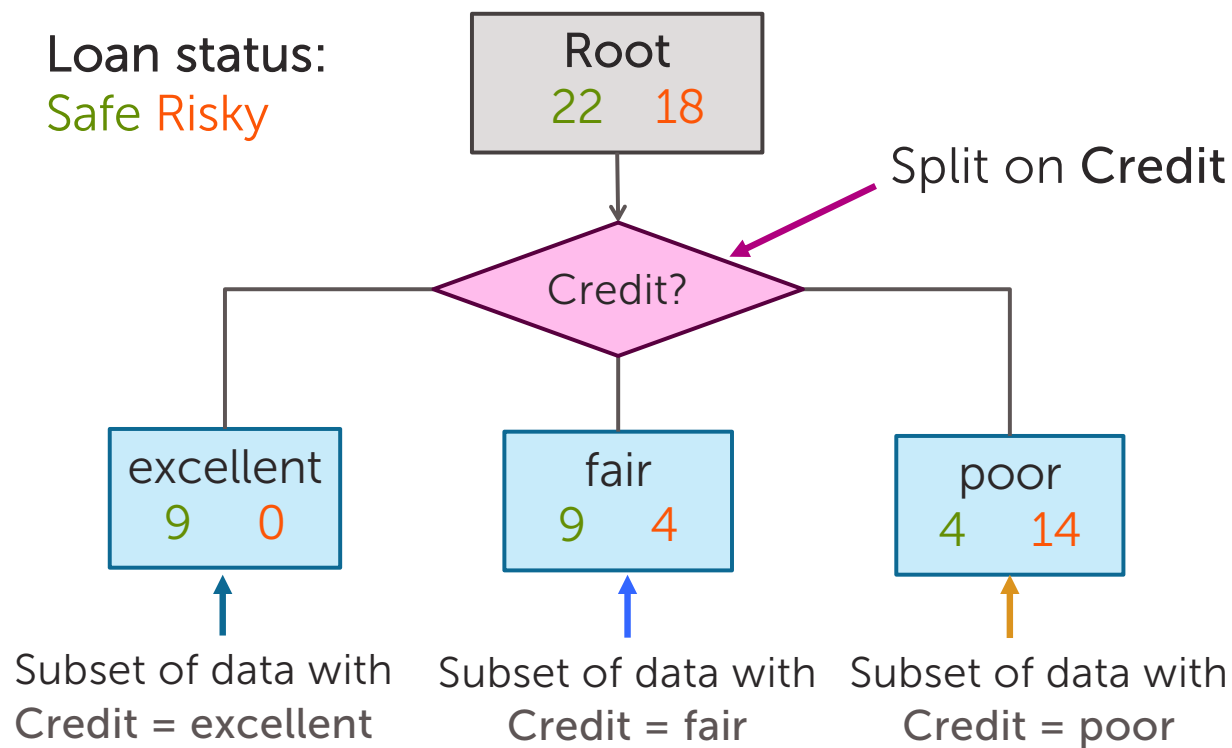


# Compact visual notation: Root node

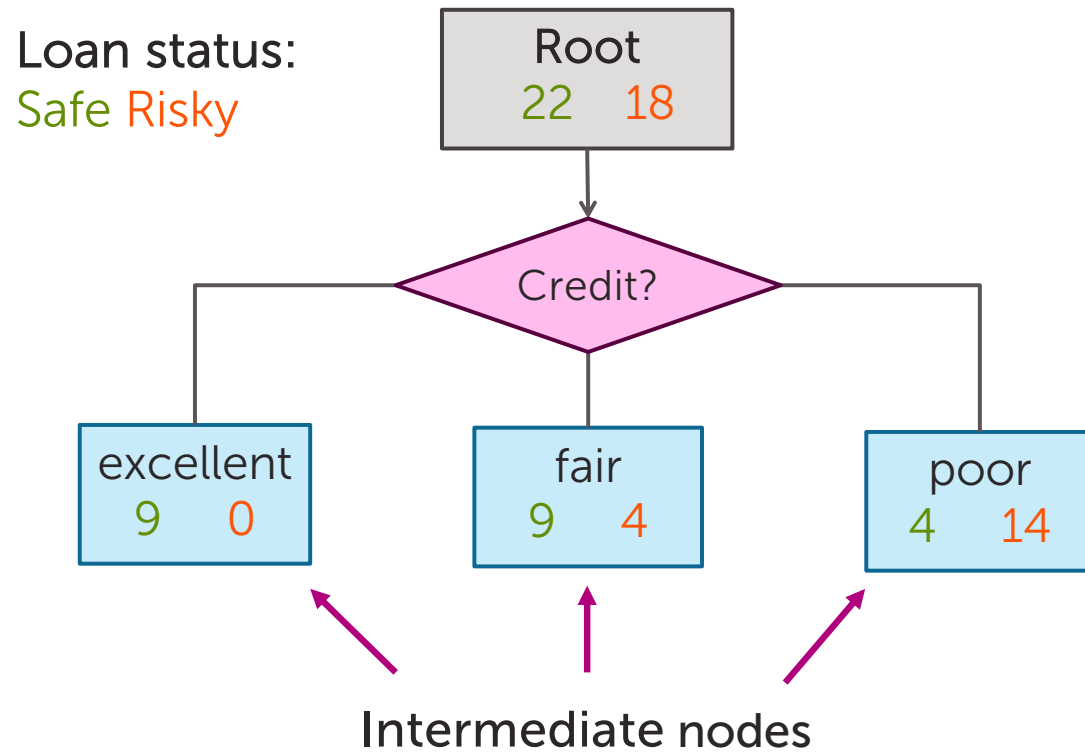
Loan status: **Safe** **Risky**



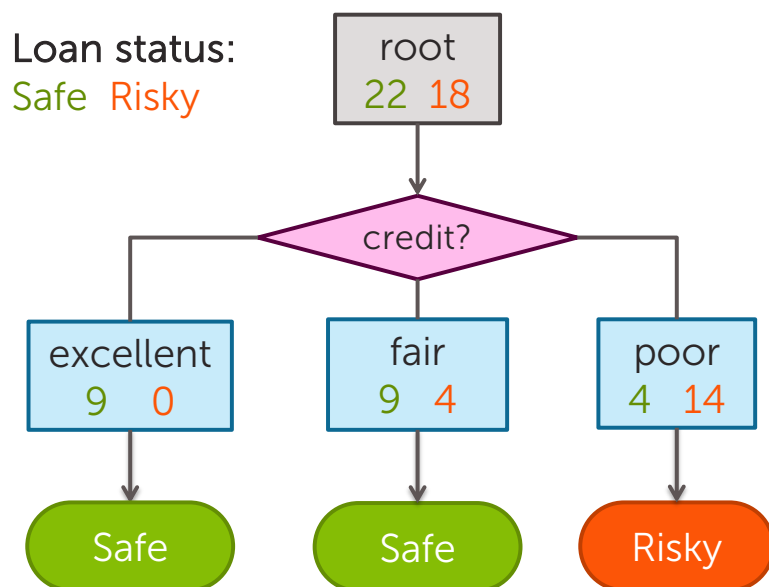
# Decision stump: Single level tree



# Visual notation: Intermediate nodes



# Making predictions with a decision stump



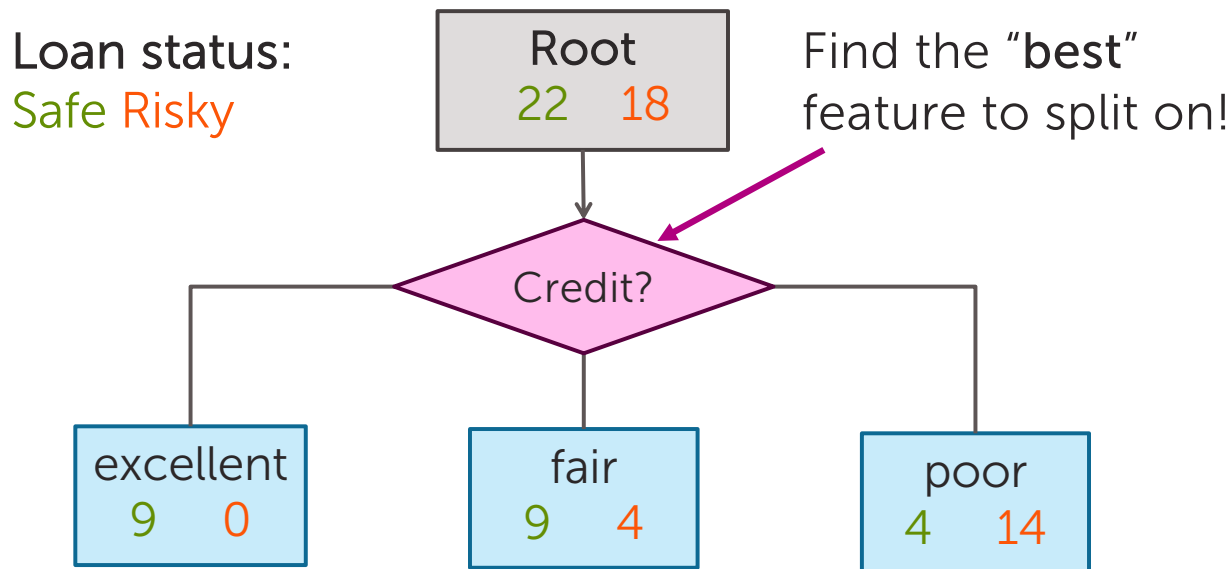
For each intermediate node,  
set  $\hat{y}$  = majority value



## Selecting best feature to split on



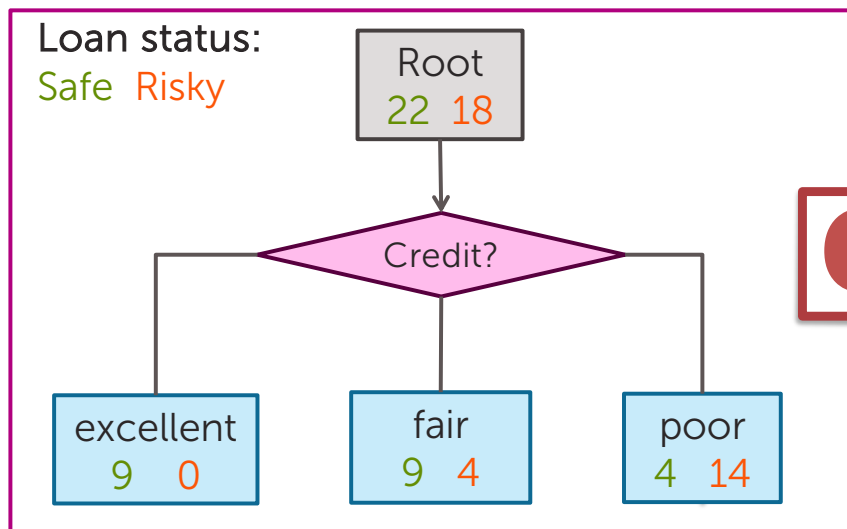
# How do we learn a decision stump?





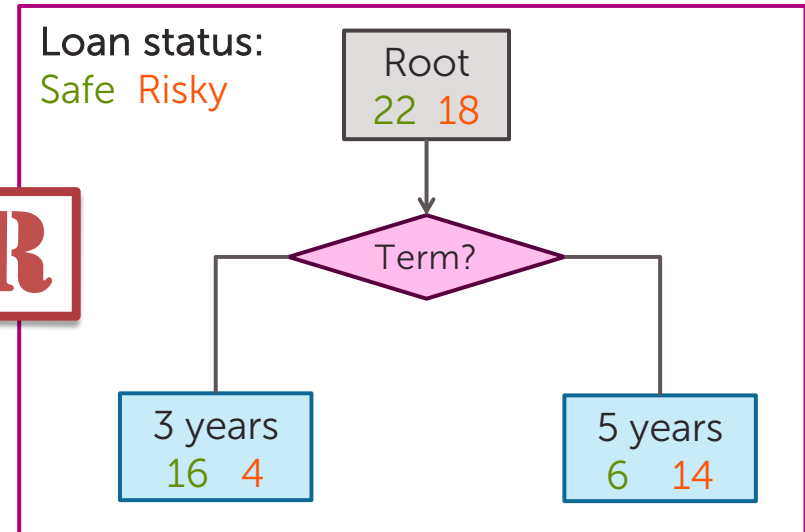
# How do we select the best feature?

## Choice 1: Split on Credit



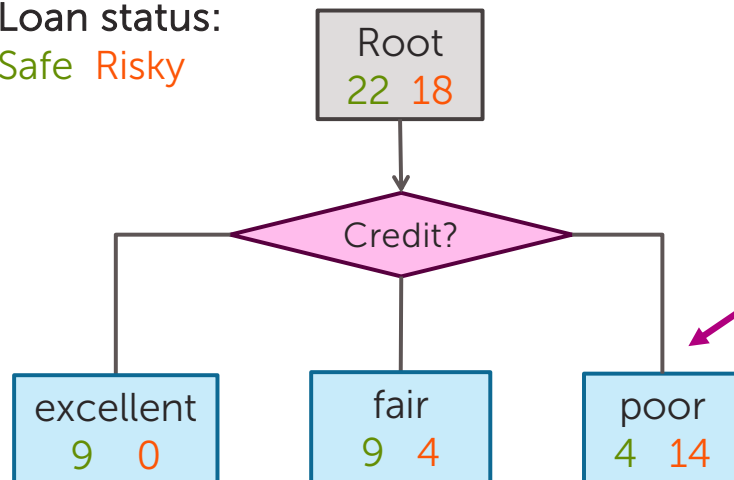
OR

## Choice 2: Split on Term



# How do we measure effectiveness of a split?

Loan status:  
Safe Risky




**Idea:** Calculate classification error of this decision stump

$$\text{Error} = \frac{\# \text{ mistakes}}{\# \text{ data points}}$$

# Calculating classification error

- **Step 1:**  $\hat{y}$  = class of majority of data in node
- **Step 2:** Calculate classification error of predicting  $\hat{y}$  for this data




Error = \_\_\_\_\_  
=

Tree	Classification error
(root)	0.45

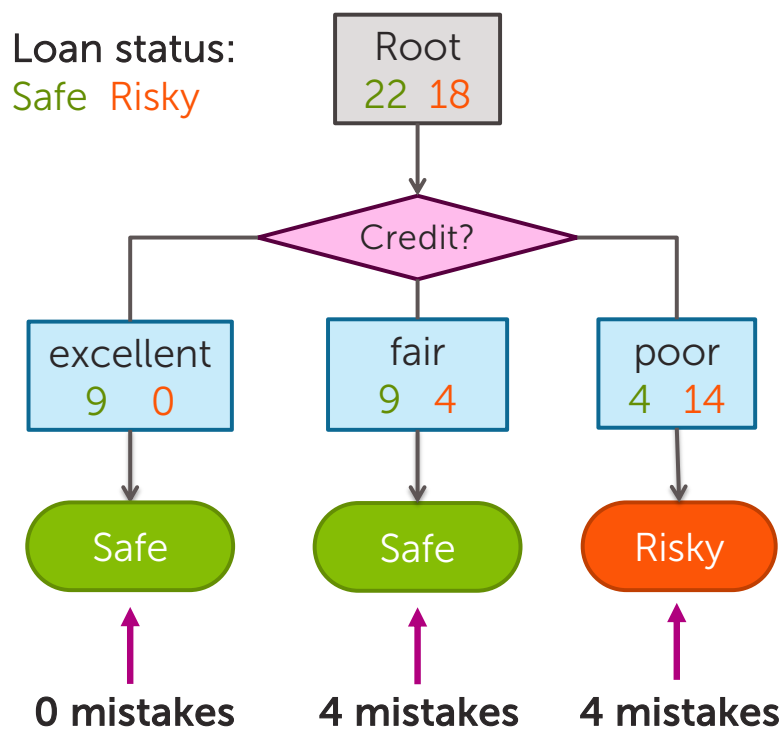
# Choice 1: Split on **Credit** history?

## Choice 1: Split on Credit



# Split on **Credit**: Classification error

## Choice 1: Split on Credit

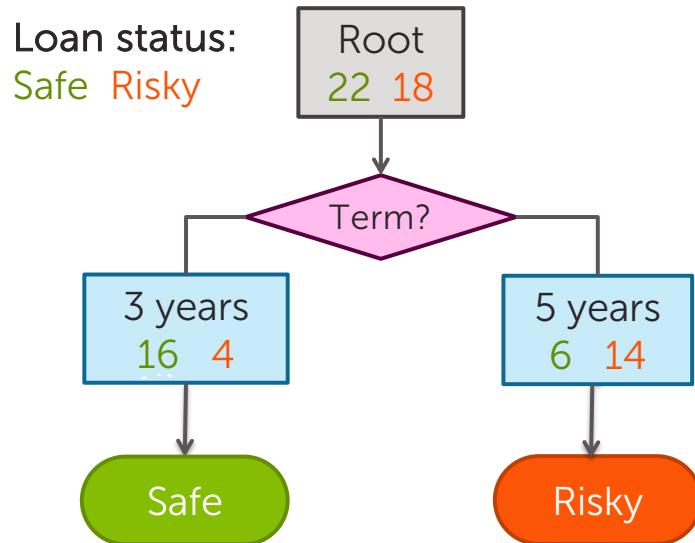


Error = \_\_\_\_\_  
=

Tree	Classification error
(root)	0.45
Split on credit	0.2

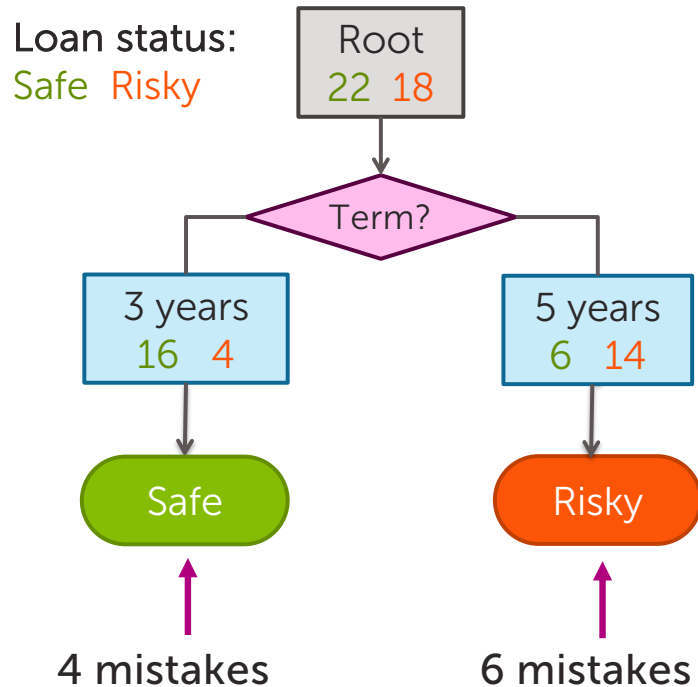
# Choice 2: Split on Term?

## Choice 2: Split on Term



# Evaluating the split on Term

## Choice 2: Split on Term




Error = \_\_\_\_\_  
=

Tree	Classification error
(root)	0.45
Split on credit	0.2
Split on term	0.25

# Choice 1 vs Choice 2: Comparing split on Credit vs Term


Tree	Classification error
(root)	0.45
split on credit	0.2
split on loan term	0.25

## Choice 1: Split on Credit



**OR**

## Choice 2: Split on Term





# Feature split selection algorithm

- Given a subset of data  $M$  (a node in a tree)
- For each feature  $h_i(x)$ :
  1. Split data of  $M$  according to feature  $h_i(x)$
  2. Compute classification error of split
- Chose feature  $h^*(x)$  with lowest classification error

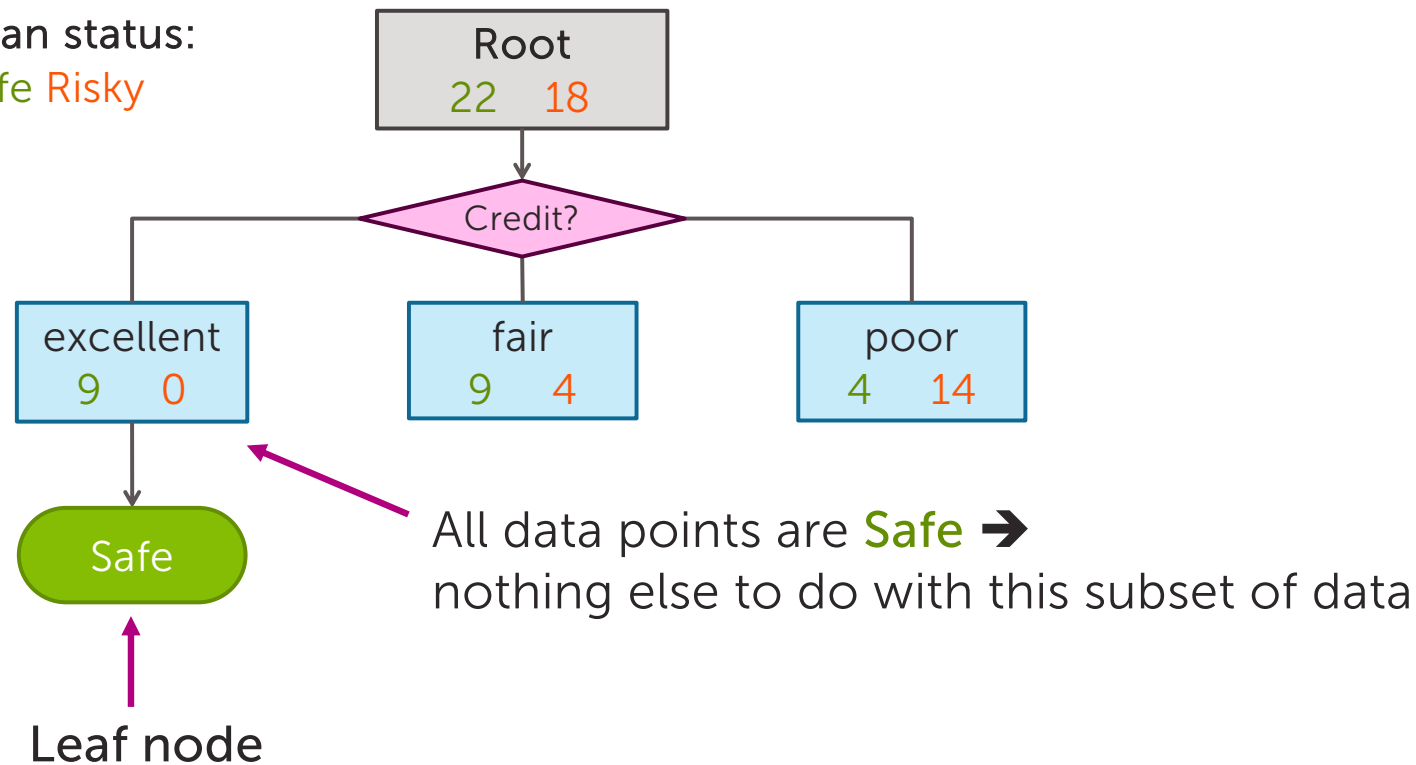


## Recursion & Stopping conditions



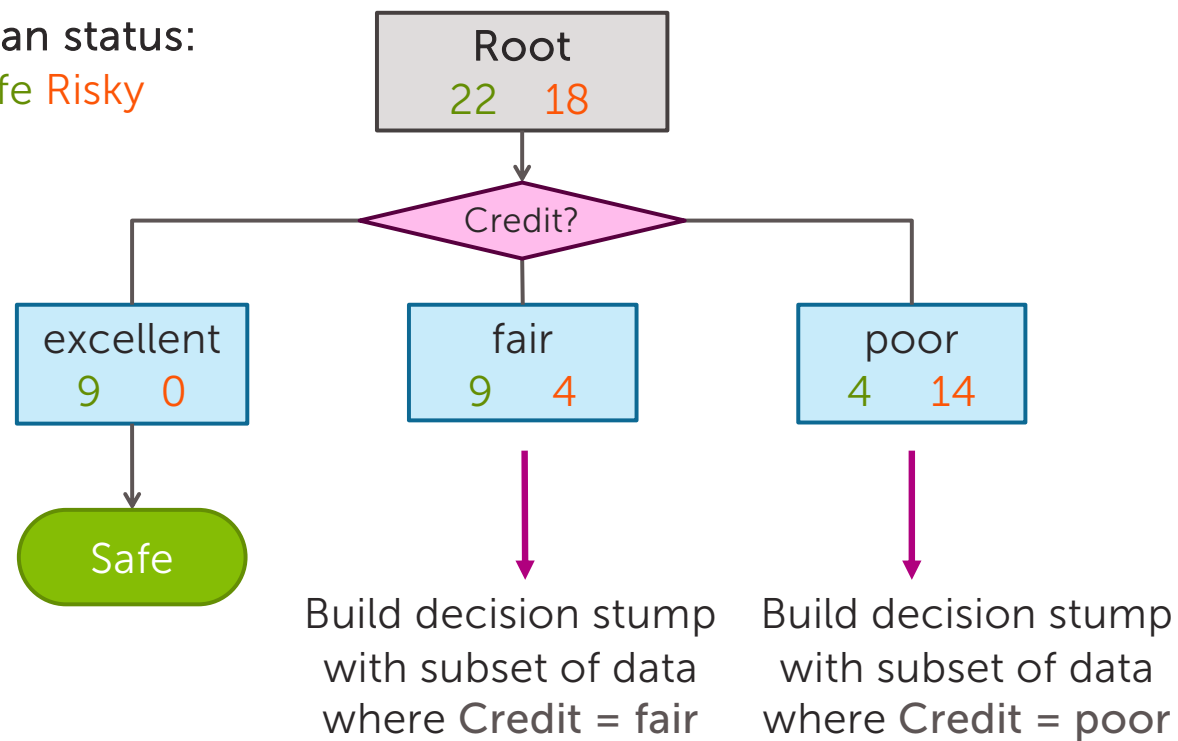
# We've learned a decision stump, what next?

Loan status:  
Safe Risky

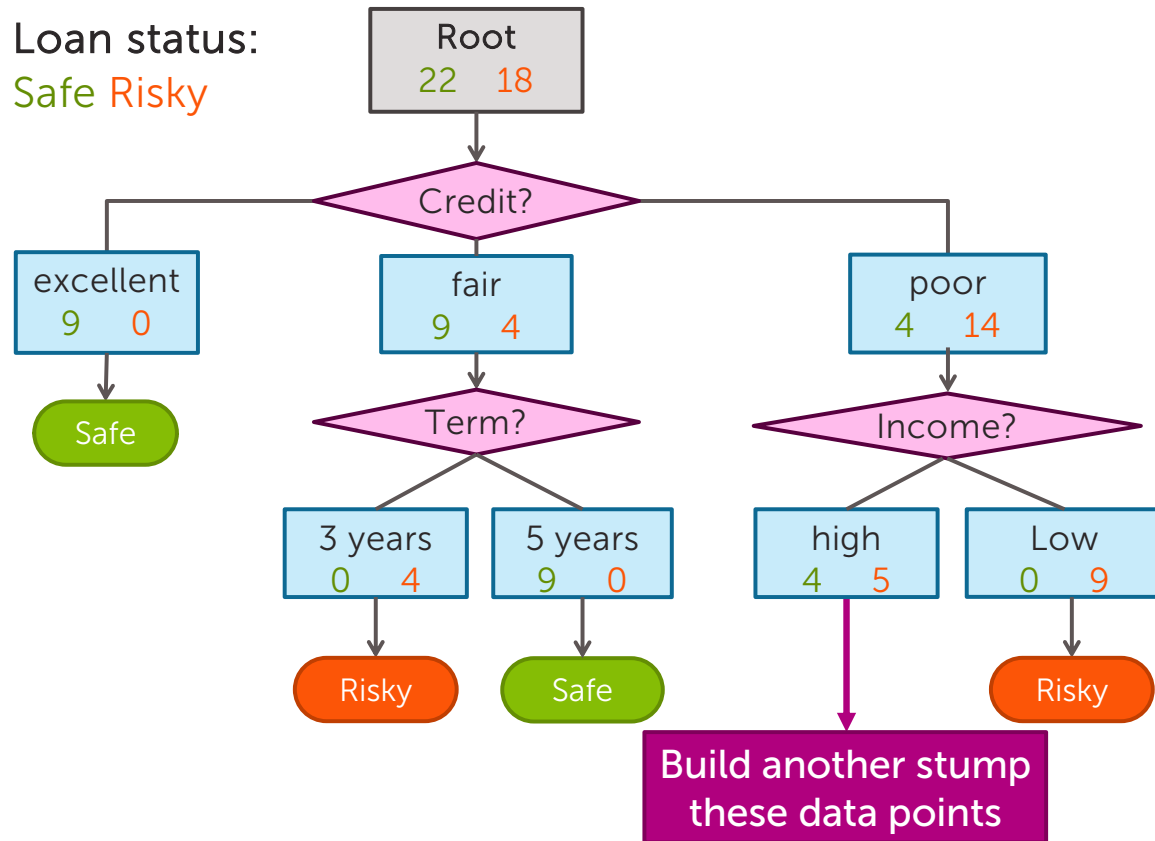


# Tree learning = Recursive stump learning

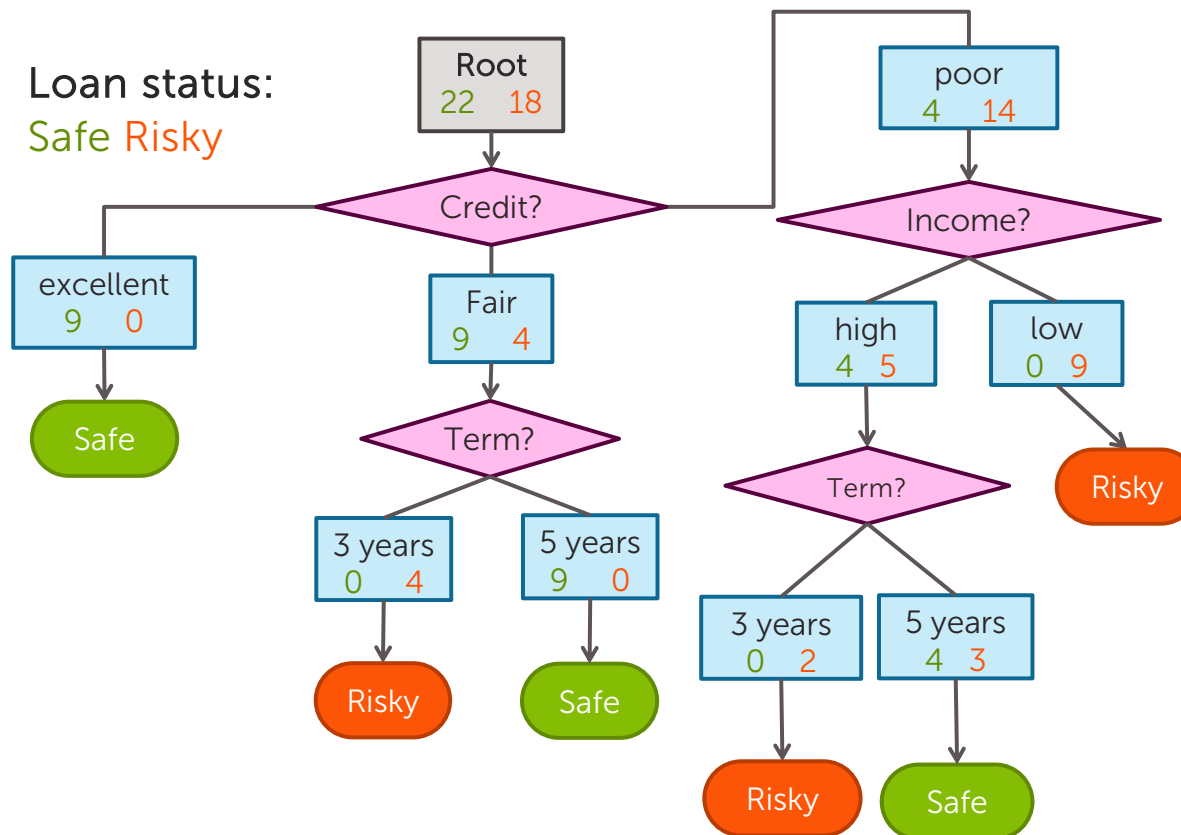
Loan status:  
Safe Risky



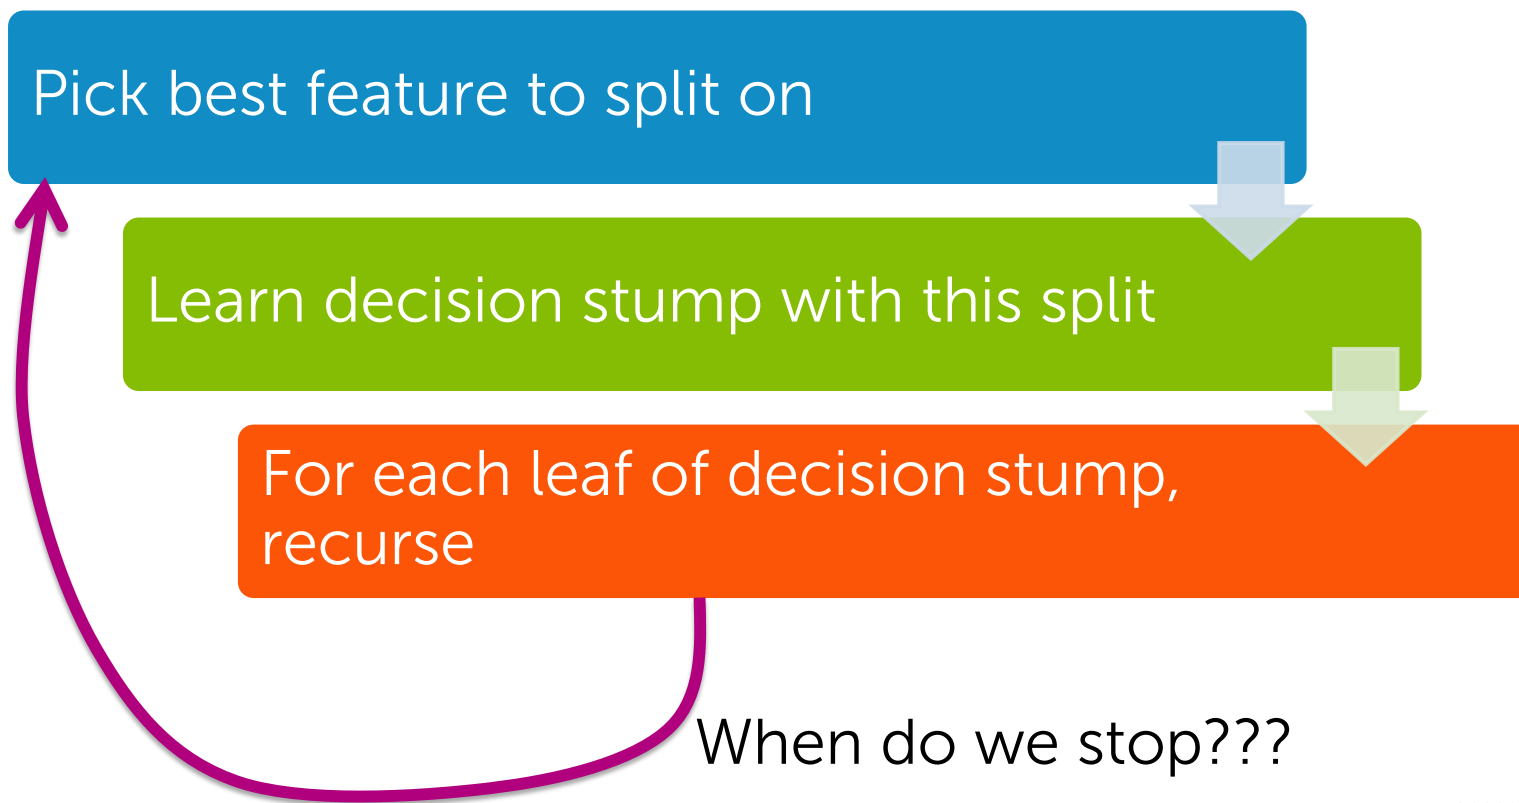
# Second level



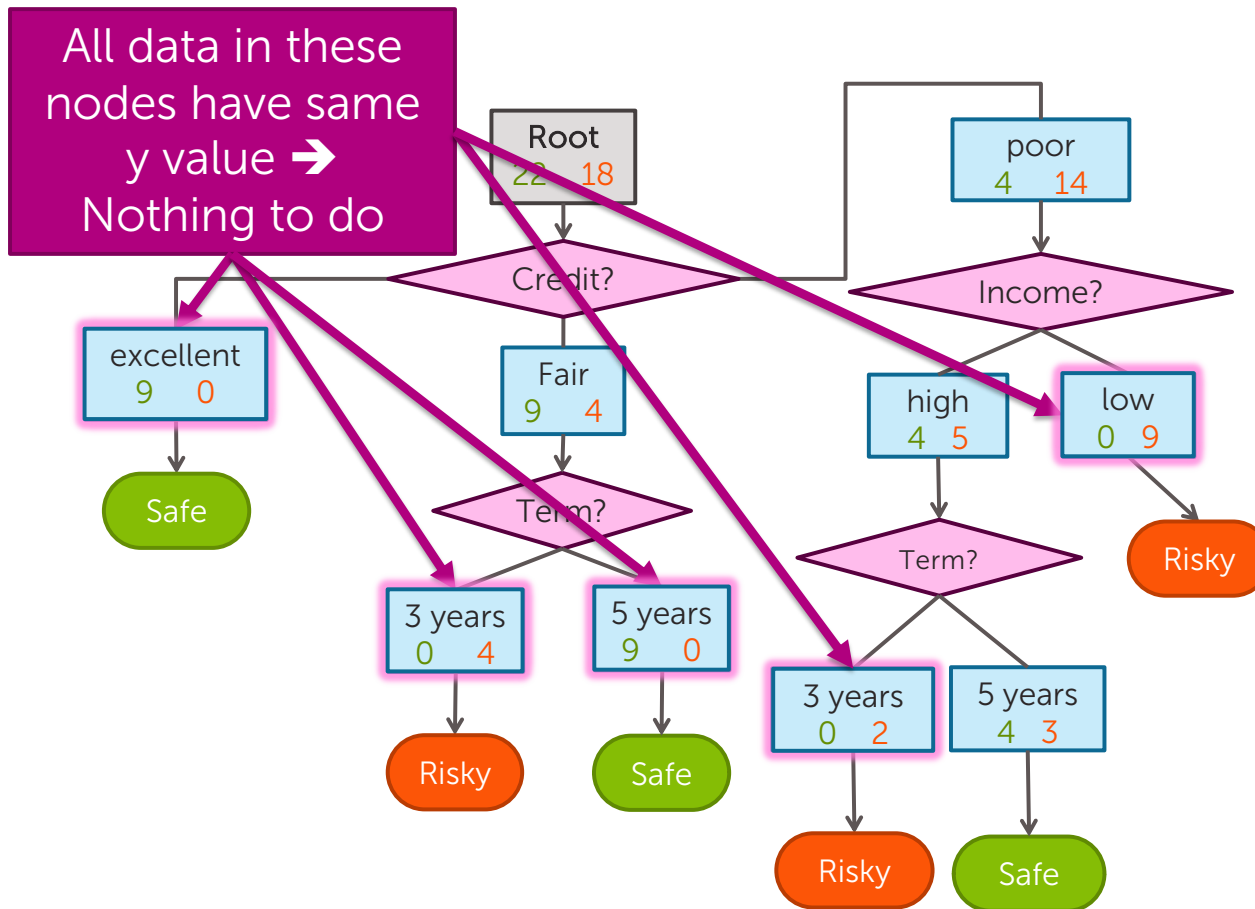
# Final decision tree



# Simple greedy decision tree learning



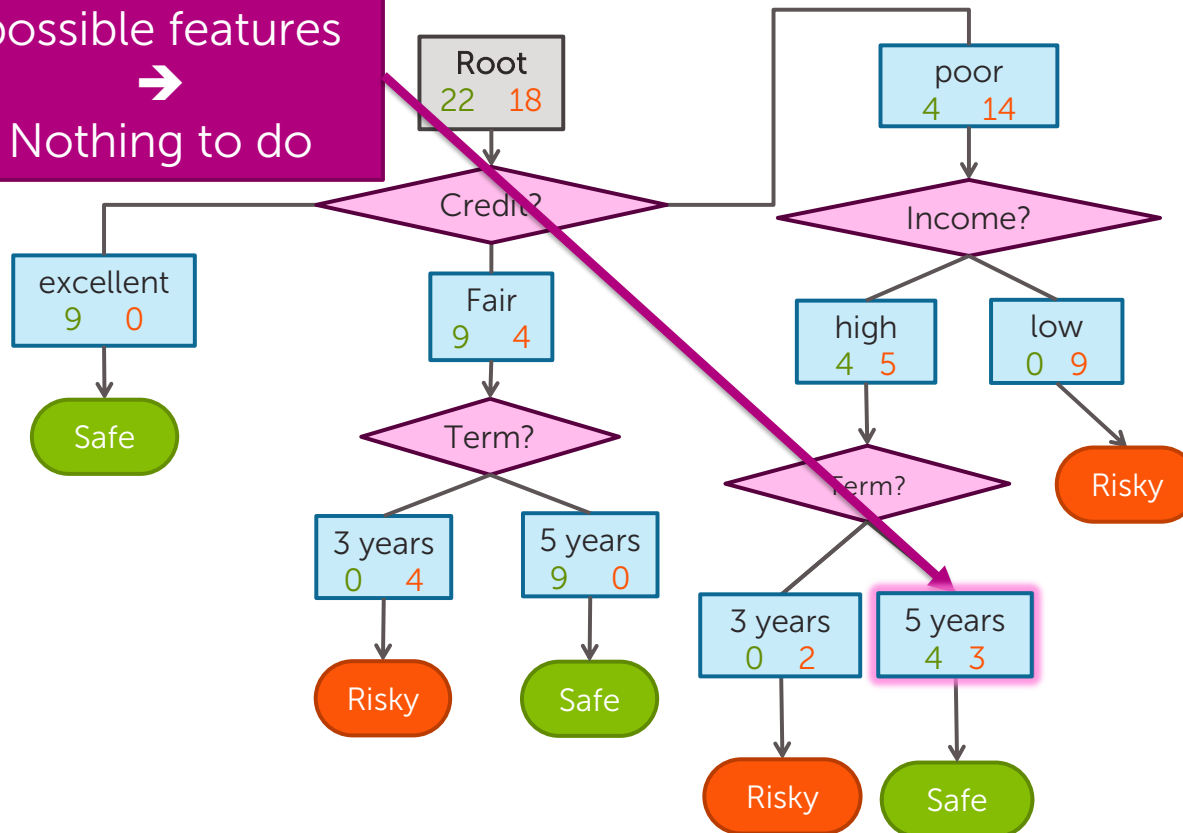
# Stopping condition 1: All data agrees on y





# Stopping condition 2: Already split on all features

Already split on all possible features  
→  
Nothing to do



# Greedy decision tree learning

- **Step 1:** Start with an empty tree

- **Step 2:** Select a feature to split data

- For each split of the tree:

- **Step 3:** If nothing more to do, make predictions

- **Step 4:** Otherwise, go to **Step 2** & continue (recurse) on this split

Pick feature split  
leading to lowest  
classification error

Stopping  
conditions

Recursion

# Is this a good idea?



Proposed stopping condition 3:  
Stop if no split reduces the  
classification error

## Stopping condition 3: Don't stop if error doesn't decrease???

$$y = \mathbf{x}[1] \text{ xor } \mathbf{x}[2]$$

$\mathbf{x}[1]$	$\mathbf{x}[2]$	$y$
False	False	False
False	True	True
True	False	True
True	True	False

y values  
True False

Root
2 2


Error = \_\_\_\_\_  
=

Tree	Classification error
(root)	0.5

# Consider split on x[1]

$$y = x[1] \text{ xor } x[2]$$

x[1]	x[2]	y
False	False	False
False	True	True
True	False	True
True	True	False



Error = \_\_\_\_\_  
=


Tree	Classification error
(root)	0.5
Split on x[1]	0.5

# Consider split on x[2]

$$y = \mathbf{x}[1] \text{ xor } \mathbf{x}[2]$$

x[1]	x[2]	y
False	False	False
False	True	True
True	False	True
True	True	False

y values  
True False



$$\text{Error} = \frac{1+1}{2+2} = 0.5$$

Neither features  
improve training error...  
Stop now???


Tree	Classification error
(root)	0.5
Split on x[1]	0.5
Split on x[2]	0.5

# Final tree with stopping condition 3

$$y = \mathbf{x}[1] \text{ xor } \mathbf{x}[2]$$

$\mathbf{x}[1]$	$\mathbf{x}[2]$	$y$
False	False	False
False	True	True
True	False	True
True	True	False

y values  
True False



Tree	Classification error
with stopping condition 3	0.5


# Without stopping condition 3

**Condition 3** (stopping when training error doesn't improve) is **not** recommended!

$$y = \mathbf{x}[1] \text{ xor } \mathbf{x}[2]$$

$\mathbf{x}[1]$	$\mathbf{x}[2]$	$y$
False	False	False
False	True	True
True	False	True
True	True	False

Tree	Classification error
with stopping condition 3	0.5
without stopping condition 3	





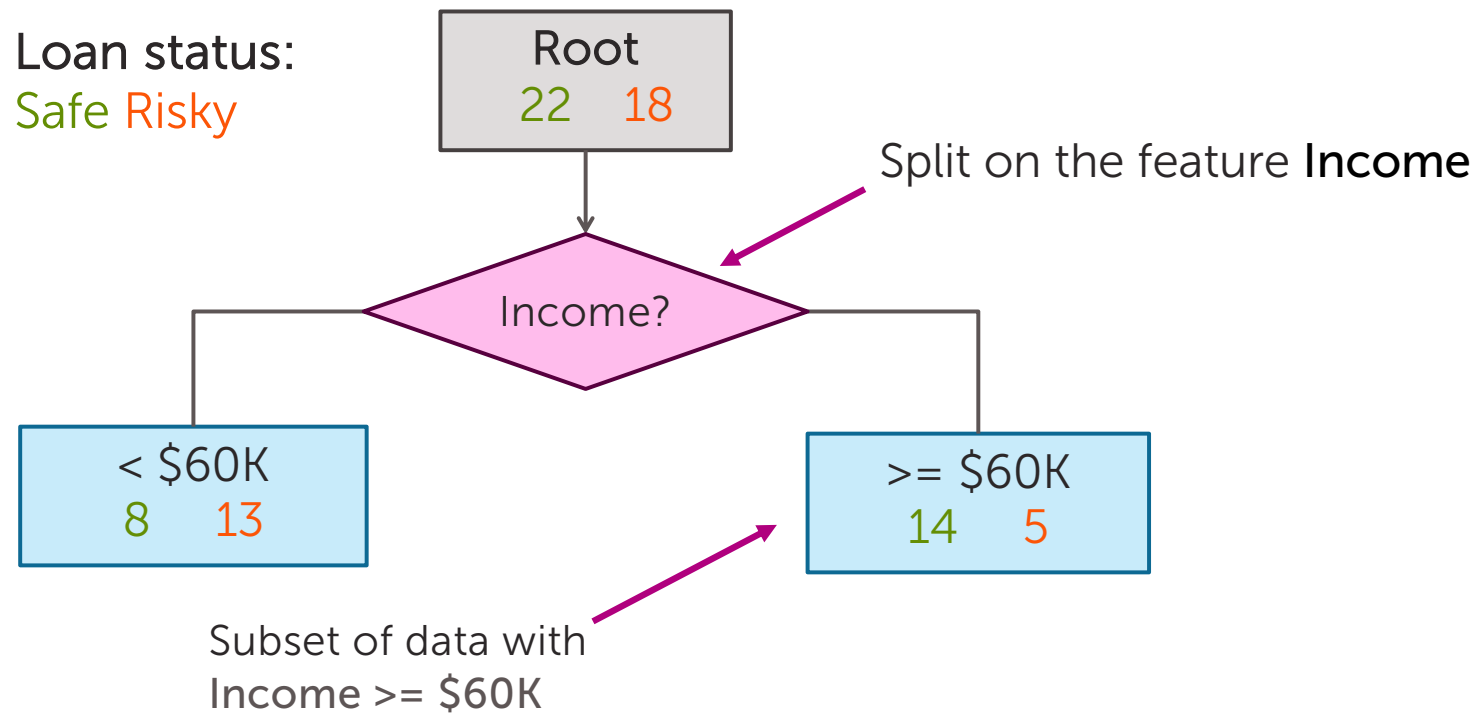
# Decision tree learning:

## *Real valued features*

# How do we use real values inputs?

Income	Credit	Term	y
\$105 K	excellent	3 yrs	Safe
\$112 K	good	5 yrs	Risky
\$73 K	fair	3 yrs	Safe
\$69 K	excellent	5 yrs	Safe
\$217 K	excellent	3 yrs	Risky
\$120 K	good	5 yrs	Safe
\$64 K	fair	3 yrs	Risky
\$340 K	excellent	5 yrs	Safe
\$60 K	good	3 yrs	Risky


# Threshold split



# Finding the best threshold split


Infinite possible values of  $t$

Income =  $t^*$




# Consider a threshold between points

Same **classification error** for any threshold split between  $v_A$  and  $v_B$



# Only need to consider mid-points


Finite number of splits to consider



# Threshold split selection algorithm


- **Step 1:** Sort the values of a feature  $h_j(\mathbf{x})$  :  
Let  $\{v_1, v_2, v_3, \dots v_N\}$  denote sorted values
- **Step 2:**
  - For  $i = 1 \dots N-1$ 
    - Consider split  $t_i = (v_i + v_{i+1}) / 2$
    - Compute classification error for threshold split  $h_j(\mathbf{x}) \geq t_i$
  - Chose the  $t^*$  with the lowest classification error

# Visualizing the threshold split







# Split on Age $\geq 38$



## Depth 2: Split on Income $\geq$ \$60K



# Each split partitions the 2-D space





# Decision trees vs logistic regression:



## *Example*

# Logistic regression

Feature	Value	Weight Learned
$h_0(\mathbf{x})$	1	0.22
$h_1(\mathbf{x})$	$x[1]$	1.12
$h_2(\mathbf{x})$	$x[2]$	-1.07




# Depth 1: Split on $x[1]$






y values

- +




# Depth 2




# Threshold split caveat

For threshold splits,  
same feature can  
be used multiple  
times







# Decision boundaries




Depth 1



Depth 2




Depth 10




# Comparing decision boundaries

## Decision Tree



---

## Logistic Regression



# Summary of decision trees

# What you can do now

- Define a decision tree classifier
- Interpret the output of a decision trees
- Learn a decision tree classifier using greedy algorithm
- Traverse a decision tree to make predictions
  - Majority class predictions
- Tackle continuous and discrete features

## Annotated decisiontrees Slides



# Decision Trees

CS229: Machine Learning

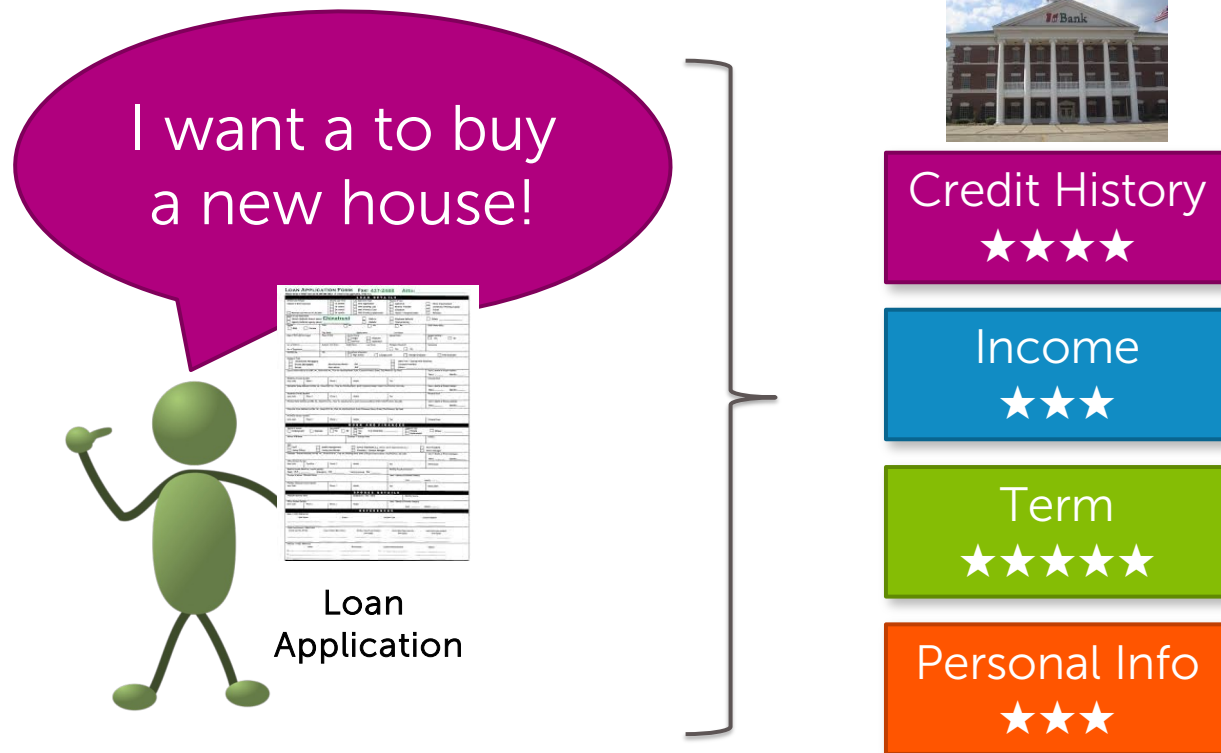
Carlos Guestrin

Stanford University

Slides include content developed by and co-developed with  
Emily Fox

# Predicting potential loan defaults

# What makes a loan risky?






# Credit history explained

Did I pay previous  
loans on time?




**Example:**  
excellent, good, or  
fair



# Income

What's my income?


**Example:**  
\$80K per year



# Loan terms

How soon do I need to pay the loan?


**Example:** 3 years,  
5 years,...




# Personal information

Age, reason for the loan, marital status,...


**Example:** Home loan for a married couple



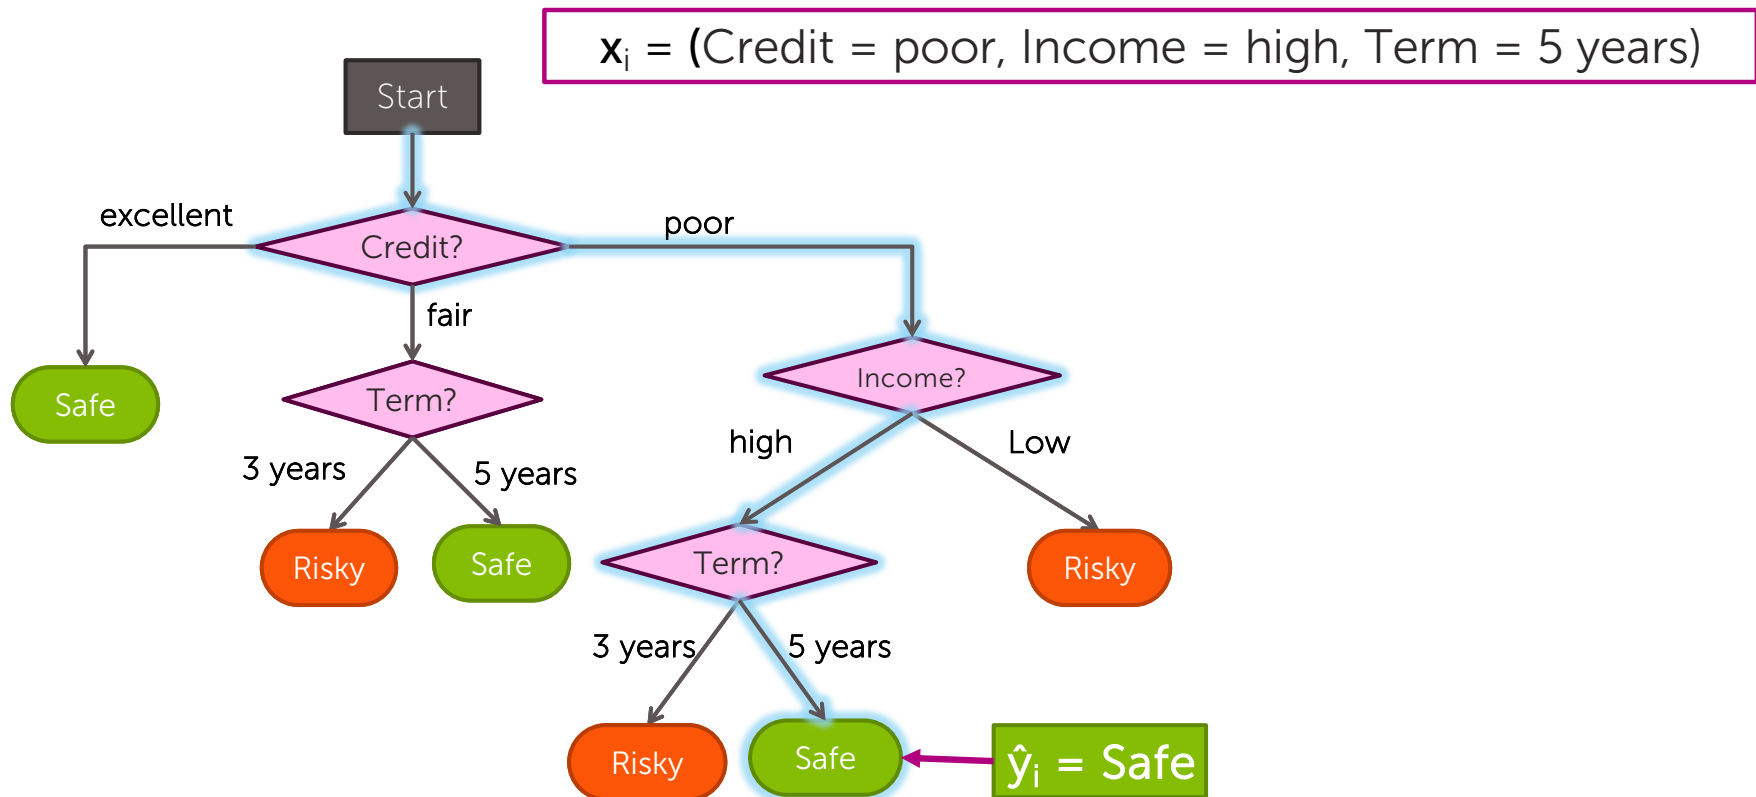
# Classifier review



# This module ... decision trees



# Scoring a loan application




# Decision tree learning task



# Decision tree learning problem

Training data:  $N$  observations  $(\mathbf{x}_i, y_i)$

Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe



# Quality metric: Classification error

- Error measures fraction of mistakes


$$\text{Error} = \frac{\text{\# incorrect predictions}}{\text{\# examples}}$$

- Best possible value : 0.0
- Worst possible value: 1.0

# How do we find the best tree?

Exponentially large number of possible trees makes decision tree learning **hard**!

Learning the smallest decision tree is an *NP-hard problem* [Hyafil & Rivest '76]



# Greedy decision tree learning


# Our training data table

Assume  $N = 40$ , 3 features

Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe


# Start with all the data

Loan status: **Safe** **Risky**




# Compact visual notation: Root node

Loan status: **Safe** **Risky**




# Decision stump: Single level tree






# Visual notation: Intermediate nodes



# Making predictions with a decision stump




For each intermediate node,  
set  $\hat{y}$  = majority value



## Selecting best feature to split on



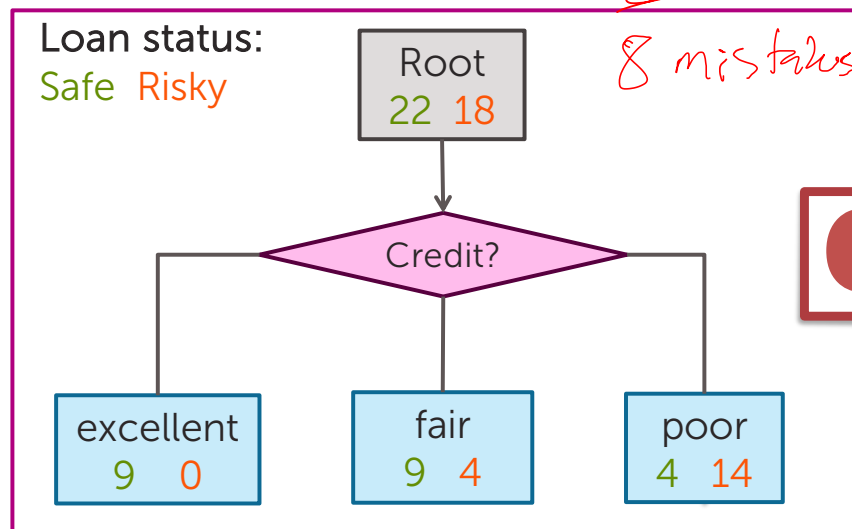
# How do we learn a decision stump?



# How do we select the best feature?

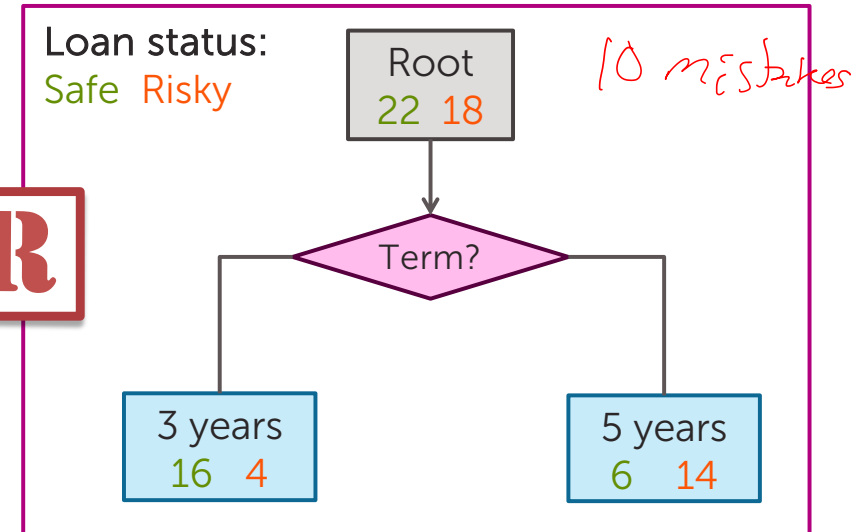
*which one makes fewer mistakes?*

## Choice 1: Split on Credit




**OR**

## Choice 2: Split on Term



# How do we measure effectiveness of a split?

Loan status:  
Safe Risky

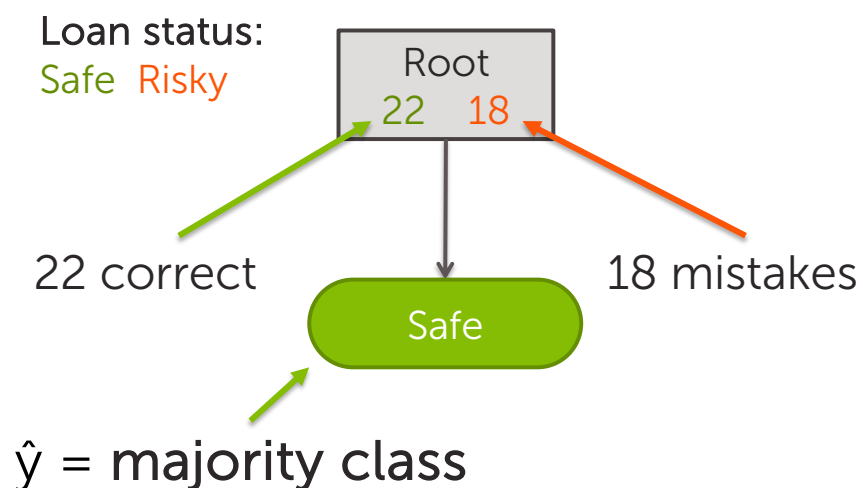


**Idea:** Calculate classification error of this decision stump

$$\text{Error} = \frac{\# \text{ mistakes}}{\# \text{ data points}}$$

# Calculating classification error

- **Step 1:**  $\hat{y}$  = class of majority of data in node
- **Step 2:** Calculate classification error of predicting  $\hat{y}$  for this data




$$\text{Error} = \frac{18}{40} = 0.45$$

Tree	Classification error
(root)	0.45

# Choice 1: Split on **Credit** history?

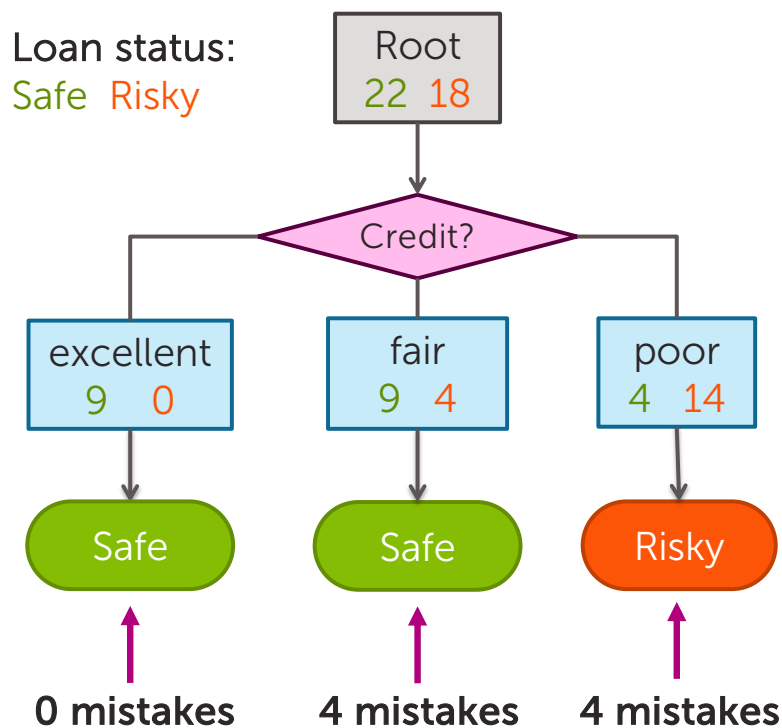
## Choice 1: Split on Credit





# Split on **Credit**: Classification error

## Choice 1: Split on Credit




$$\text{Error} = \frac{8}{40} = 0.2$$

Tree	Classification error
(root)	0.45
Split on credit	0.2


# Choice 2: Split on Term?

## Choice 2: Split on Term



# Evaluating the split on Term

## Choice 2: Split on Term



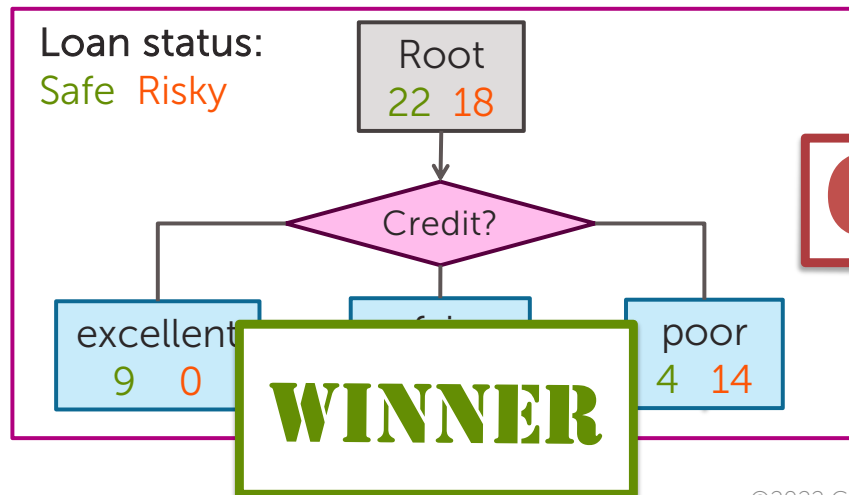
$$\text{Error} = \frac{10}{40} = 0.25$$

Tree	Classification error
(root)	0.45
Split on credit	0.2
Split on term	0.25

# Choice 1 vs Choice 2: Comparing split on Credit vs Term


Tree	Classification error
(root)	0.45
split on credit	0.2
split on loan term	0.25

## Choice 1: Split on Credit



**OR**

## Choice 2: Split on Term



# Feature split selection algorithm

- Given a subset of data  $M$  (a node in a tree)
- For each feature  $h_i(x)$ :
  1. Split data of  $M$  according to feature  $h_i(x)$
  2. Compute classification error of split
- Chose feature  $h^*(x)$  with lowest classification error



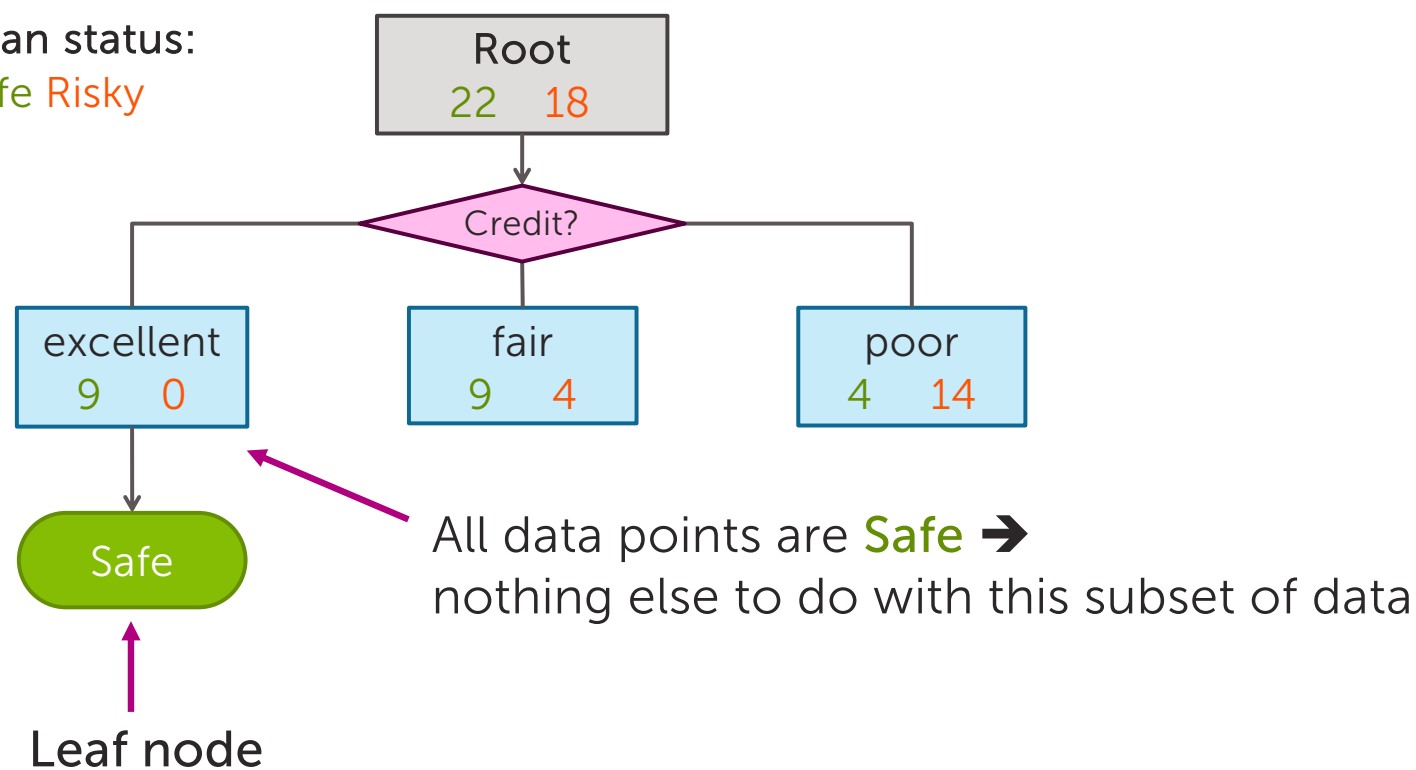
## Recursion & Stopping conditions



# We've learned a decision stump, what next?


Loan status:

Safe Risky




# Tree learning = Recursive stump learning

Loan status:  
Safe Risky







# Second level



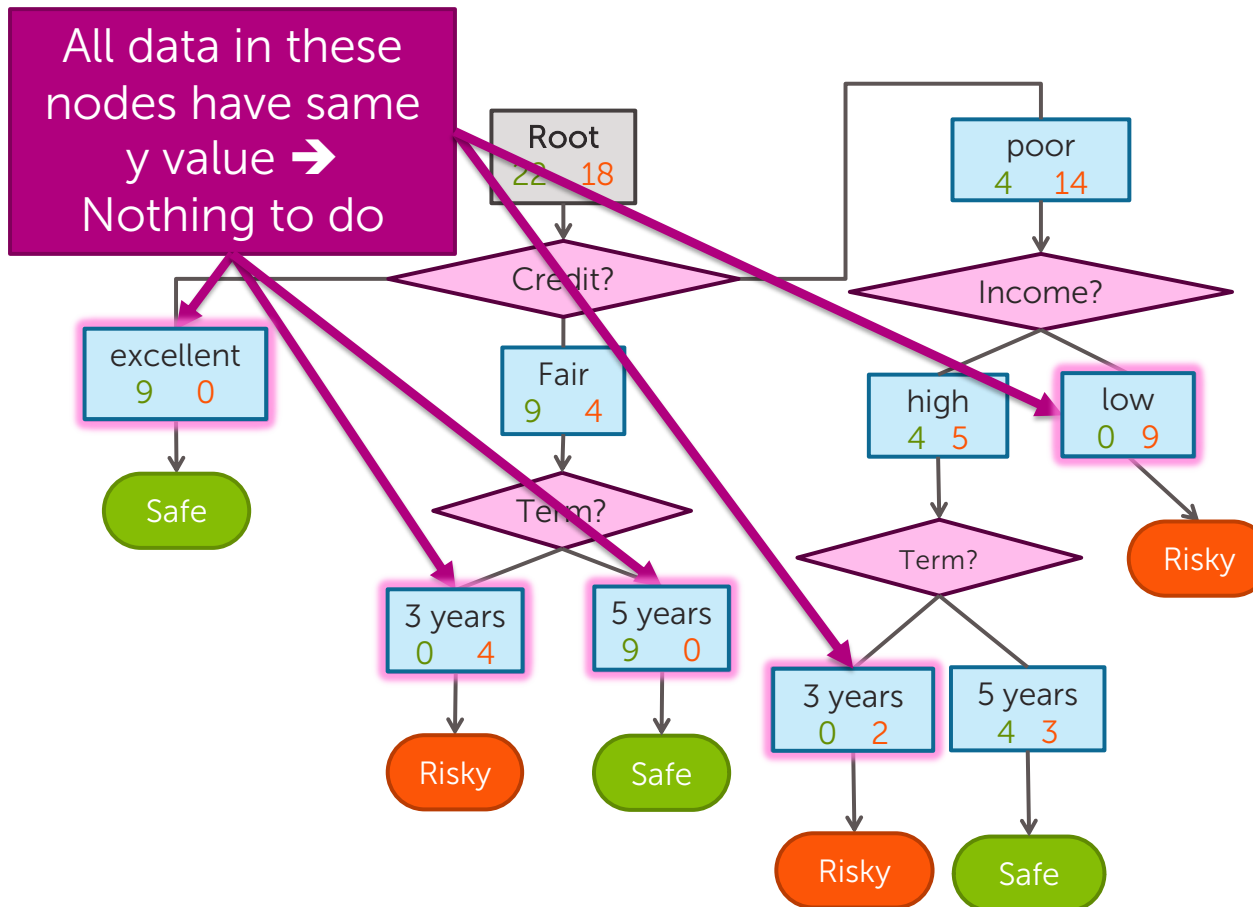
# Final decision tree



# Simple greedy decision tree learning

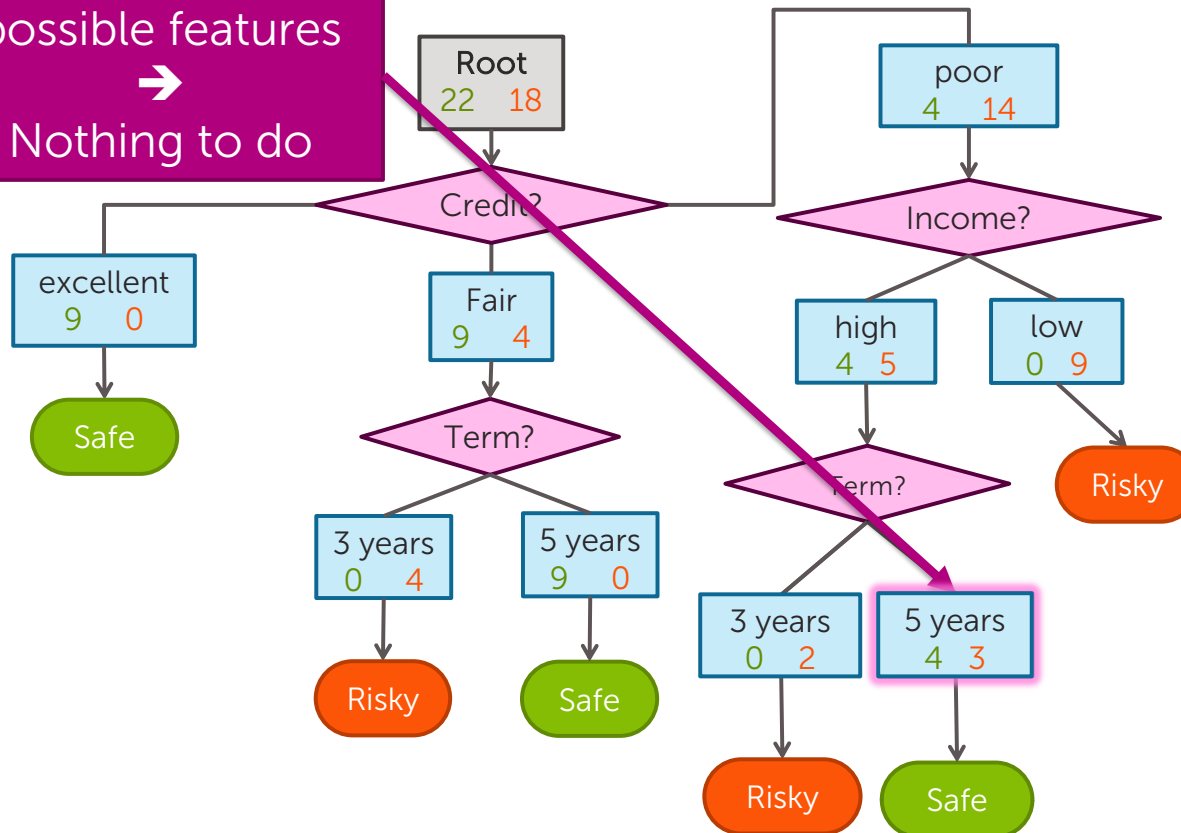


# Stopping condition 1: All data agrees on y



# Stopping condition 2: Already split on all features

Already split on all  
possible features  
→  
Nothing to do



# Greedy decision tree learning

- **Step 1:** Start with an empty tree

- **Step 2:** Select a feature to split data

- For each split of the tree:

- **Step 3:** If nothing more to do, make predictions

- **Step 4:** Otherwise, go to **Step 2** & continue (recurse) on this split

Pick feature split leading to lowest classification error

Stopping conditions


Recursion

# Is this a good idea?



Proposed stopping condition 3:  
Stop if no split reduces the  
classification error

# Stopping condition 3: Don't stop if error doesn't decrease???



$$y = \mathbf{x}[1] \text{ xor } \mathbf{x}[2]$$

$\mathbf{x}[1]$	$\mathbf{x}[2]$	$y$
False	False	False
False	True	True
True	False	True
True	True	False

y values  
True False

Root
2 2

$$\text{Error} = \frac{2}{4} = 0.5$$


Tree	Classification error
(root)	0.5



# Consider split on x[1]

$$y = x[1] \text{ xor } x[2]$$

x[1]	x[2]	y
False	False	False
False	True	True
True	False	True
True	True	False



$$\text{Error} = \frac{2}{4} = 0.5$$


Tree	Classification error
(root)	0.5
Split on x[1]	0.5

# Consider split on x[2]

$$y = x[1] \text{ xor } x[2]$$

x[1]	x[2]	y
False	False	False
False	True	True
True	False	True
True	True	False

y values  
True False



$$\text{Error} = \frac{1+1}{2+2} = 0.5$$

Neither features  
improve training error...  
Stop now???


Tree	Classification error
(root)	0.5
Split on x[1]	0.5
Split on x[2]	0.5

# Final tree with stopping condition 3

$$y = \mathbf{x}[1] \text{ xor } \mathbf{x}[2]$$

$\mathbf{x}[1]$	$\mathbf{x}[2]$	$y$
False	False	False
False	True	True
True	False	True
True	True	False

y values  
True False




Tree	Classification error
with stopping condition 3	0.5

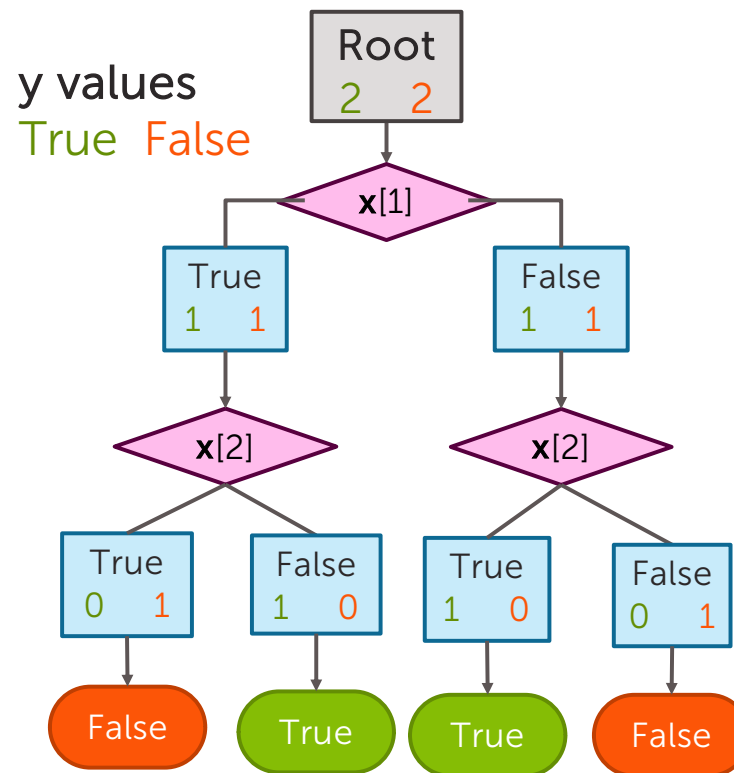
# Without stopping condition 3

**Condition 3** (stopping when training error doesn't improve) is **not** recommended!

$$y = \mathbf{x}[1] \text{ xor } \mathbf{x}[2]$$

$\mathbf{x}[1]$	$\mathbf{x}[2]$	$y$
False	False	False
False	True	True
True	False	True
True	True	False

Tree	Classification error
with stopping condition 3	0.5
without stopping condition 3	




# Decision tree learning:

## *Real valued features*

# How do we use real values inputs?

Income	Credit	Term	y
\$105 K	excellent	3 yrs	Safe
\$112 K	good	5 yrs	Risky
\$73 K	fair	3 yrs	Safe
\$69 K	excellent	5 yrs	Safe
\$217 K	excellent	3 yrs	Risky
\$120 K	good	5 yrs	Safe
\$64 K	fair	3 yrs	Risky
\$340 K	excellent	5 yrs	Safe
\$60 K	good	3 yrs	Risky


# Threshold split *→ turn continuous var into binary*



# Finding the best threshold split

Infinite possible values of  $t$


Income =  $t^*$





# Consider a threshold between points


Same **classification error** for any threshold split between  $v_A$  and  $v_B$



# Only need to consider mid-points

*Sorted data:*


Finite number of  
splits to consider




# Threshold split selection algorithm

- **Step 1:** Sort the values of a feature  $h_j(\mathbf{x})$  :  
Let  $\{v_1, v_2, v_3, \dots v_N\}$  denote sorted values
- **Step 2:**
  - For  $i = 1 \dots N-1$ 
    - Consider split  $t_i = (v_i + v_{i+1}) / 2$
    - Compute classification error for threshold split  $h_j(\mathbf{x}) \geq t_i$
  - Chose the  $t^*$  with the lowest classification error


# Visualizing the threshold split




# Split on Age $\geq 38$



## Depth 2: Split on Income $\geq$ \$60K



# Each split partitions the 2-D space





# Decision trees vs logistic regression:

## *Example*





# Logistic regression

Feature	Value	Weight Learned
$h_0(\mathbf{x})$	1	0.22
$h_1(\mathbf{x})$	$x[1]$	1.12
$h_2(\mathbf{x})$	$x[2]$	-1.07




# Depth 1: Split on $x[1]$

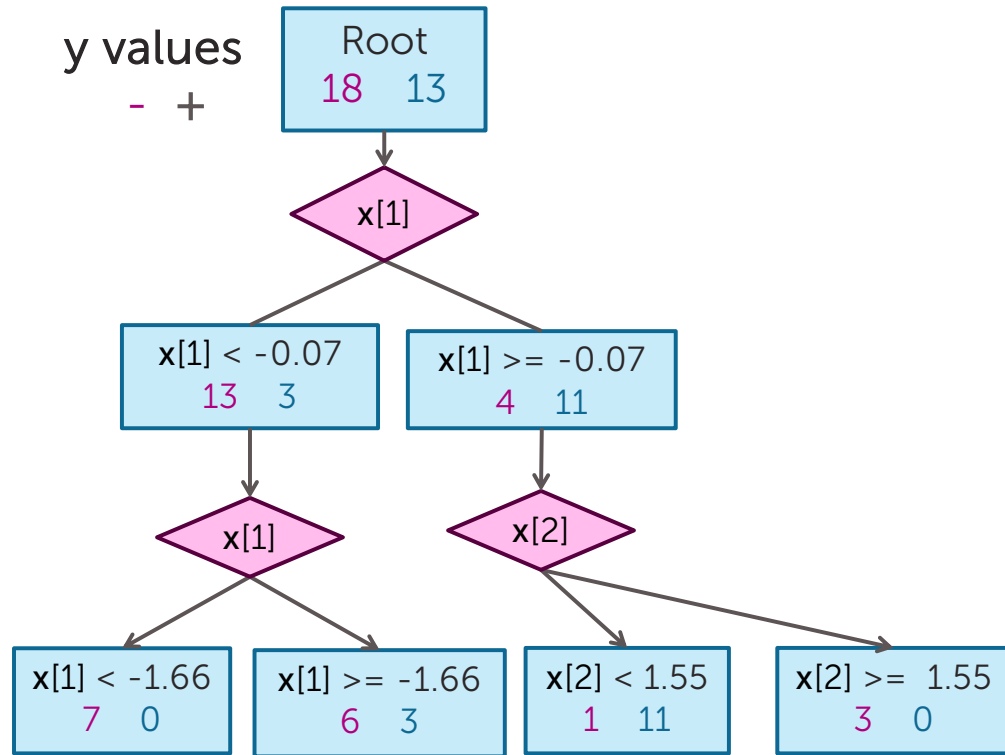




y values

- +

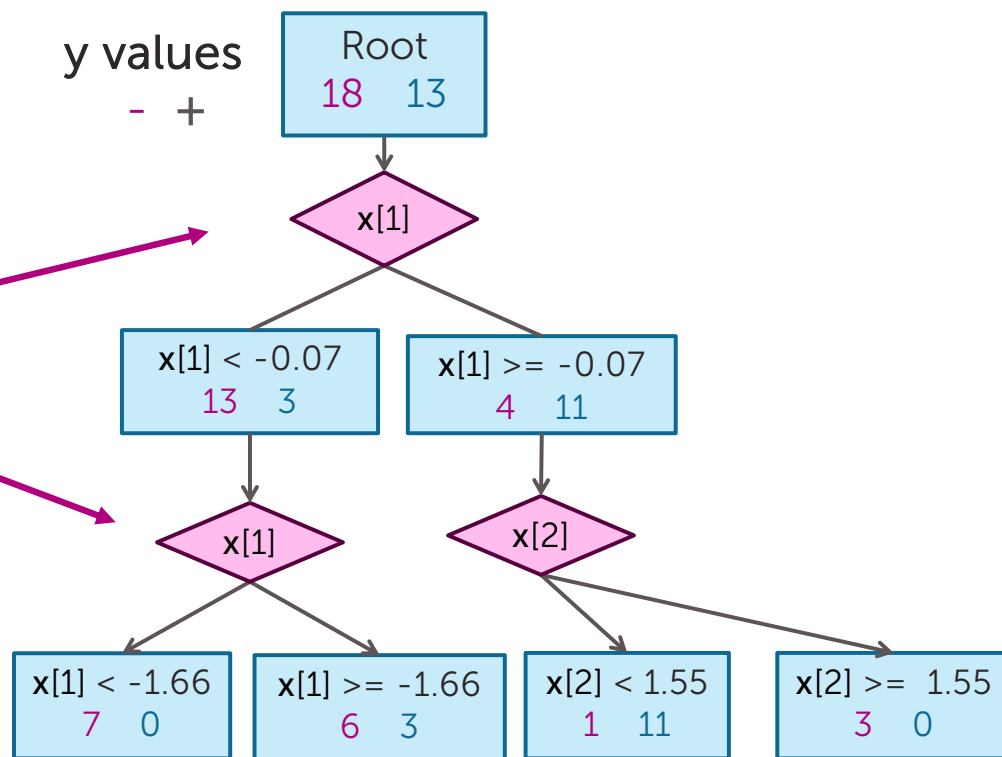


# Depth 2




# Threshold split caveat


For threshold splits,  
same feature can  
be used multiple  
times




# Decision boundaries




Depth 1



Depth 2




Depth 10




# Comparing decision boundaries

## Decision Tree



---

## Logistic Regression



# Summary of decision trees

# What you can do now

- Define a decision tree classifier
- Interpret the output of a decision trees
- Learn a decision tree classifier using greedy algorithm
- Traverse a decision tree to make predictions
  - Majority class predictions
- Tackle continuous and discrete features